

Brainwashing session : How to make T<sub>E</sub>X docs  
with tons of graphs (part II)?



# SQLite $\text{\TeX}$

## Introduction

## Basics

## Walking through a real example

Misc.

How you get started (if brainwashed)?

## Introduction

## Basics

## Walking through a real example

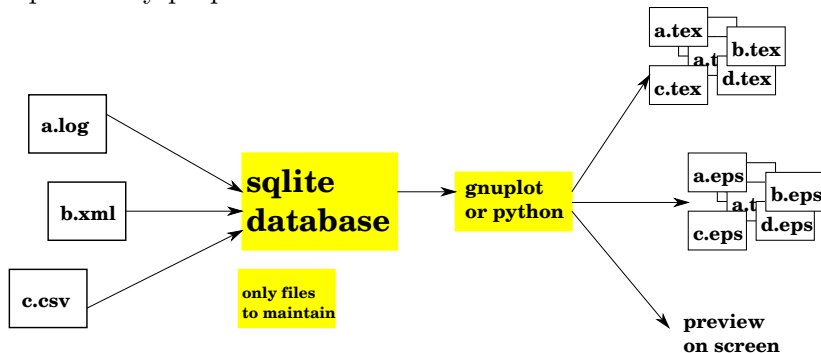
Misc.

# Context


- ▶ You want to create a T<sub>E</sub>X document with lots of graphs easily
- ▶ You want to *automate the entire process*, from running experiments to producing T<sub>E</sub>X document with graphs, so you can *painlessly repeat the experiment* and update data in the document
- ▶ With ad-hoc solutions, your directory easily screws up with too many data files and scripts you never understand a week later

## The practice

I previously proposed:



## The practice



- ▶ You maintain only 2 files for all graphs
  - ▶ an `sqlite3` database that has all raw data
  - ▶ a `gnuplot` or a script file that generates data to plot (by `gnuplot`)
- ▶ I talked about a command, `txt2sql`, which makes it straightforward to convert text files (log) to `sqlite3` database
- ▶ *Generating lots of graphs from a database was still painful, and I now address it*








## A simplest example


```
1 import smart_gnuplotter
2 g = smart_gnuplotter.smart_gnuplotter()
3 g.graphs("x*x")
```



## Writing multiple plots in a single graph


```
1 g.graphs("%(a)s * x", a=[1,2,3])
```

- The basic form is to parameterize the expression with *%(var)s* and supply its values with *var=list*




## More plots in a single graph ...

```
1 g.graphs("%(a)s*x*x+%(b)s*x+%(c)", a=[1,2,3], b=[4,5,6], c=[7,8,9])
```




## Writing multiple graphs in a single shot (1)


```
1 g.graphs("%(a)s*x*x+%(b)s*x+%(c)", a=[1,2,3], b=[4,5,6], c=[7,8,9],
2      graph_vars=["a"])
```



## Writing multiple graphs in a single shot (2)



## Writing multiple graphs in a single shot (3)



## Changing terminals and writing to files

```

1 g.graphs("x*x",
2         terminal="epslatex size 10cm,5cm",
3         output="xx")

```

- ▶ will use epslatex terminal and generate output to **xx.tex**
- ▶ for some commonly used terminal types, extensions are automatically attached, so you don't have to change output names every time you change terminal types

## Changing default terminals

Alternatively, you may set the default terminal.

```
1 g.default_terminal = "epslatex size 10cm,5cm"
2 g.graphs("x*x", output="xx")
```





- ```
1 g.default_pause=x
```




## Specifying various attributes

```

1 g.graphs("(%a)s*x*x+% (b)s*x+% (c)s", a=[1,2,3], b=[4,5,6], c=[7,8,9],
2       graph_vars=["a"], graph_attr=r'''
3 set title "$%(a)sx^2+bx+c$"
4 ''',
5       plot_attr=r'''title "$%(a)sx^2+% (b)sx+% (c)s$"''')

```



## Specifying various attributes

- ▶ `graph_attr` modifies an entire graph; it is whatever comes before the 'plot' command
- ▶ `plot_attr` modifies each plot in a graph; it is whatever comes after each plot

```
1 g.graphs(E,
2     graph_attr=graph_attr,
3     plot_attr=plot_attr)
```

≈

```
1 graph_attr
2 plot E plot_attr, E plot_attr, ...
```

## Shortcuts for frequently used attributes


Frequently used attributes can be directly set by keyword arguments. e.g.,

```

1 g.graphs("%(a)s * x * x + %(b)s * x + %(c)s",
2         a=[1,2,3], b=[4,5,6], c=[7,8,9],
3         output="graphs/mod_%(a)s_xx_bx_c",
4         graph_title="$%(a)sx^2+bx+c$",
5         graph_attr=r''set key left
6     ''',
7         plot_title="$%(a)sx^2+%(b)sx+%(c)s$",
8         plot_with="lines linewidth 2",
9         graph_vars=["a"])
```

## Shortcuts for frequently used attributes

Then you get:



## Supported graph attributes

| name        | description       | example                  |
|-------------|-------------------|--------------------------|
| terminal    | terminal used     | "epslatex size 10cm,5cm" |
| output      | file(*) to output | "my_graph_%(a)s"         |
| graph_title | graph title       | "speedup_%(a)s"          |
| xrange      | xrange            | "[0:10]"                 |
| yrange      | yrange            | "[0:10]"                 |
| boxwidth    | set boxwidth(**)  | "0.9 relative"           |

- ▶ \* : it may be extended based on terminal type.
- ▶ \*\* : effective only when you plot with boxes

The effect  $\approx$

```
1 set attr value
```



## Supported plot attributes

| name               | description                                  | example          |
|--------------------|----------------------------------------------|------------------|
| plot_title         | plot title                                   | "%(a)s"          |
| plot_with<br>using | style to plot plots<br>columns to plot plots | "boxes"<br>"1:2" |

The effect  $\approx$

```
1 plot ... using using with plot_with title "plot_title"
```

## Other things to plot

- ▶ As it is normal in gnuplot, you may plot

- ▶ datafile

```
1 g.graphs('filename', ...)
```

- ▶ output of a command

```
1 g.graphs('>cmd', ...)
```

- ▶ Besides, you may plot

- ▶ data in python list


```
1 g.graphs([(1,2),(3,4),...], ...)
```

- ▶ output of an SQLite query (later)

## Bonus : writing graphs with symbolic $x$ -axis

- ▶ With only gnuplot, it's tedious to show a graph with symbolic  $x$ -axis
- ▶ smart\_gnuplotter does all the work you need to do for it

```
1 g.graphs([("a",10), ("b",20), ("c", 30)],
2         xrange="0:", boxwidth="0.9 relative",
3         plot_with="boxes fs pattern 1")
```






## A basic example with sqlite3

Say a database example.sqlite contains the following table

```
1 select * from t
2 0|1
3 2|3
4 4|5
```

```
1 g.graphs(("example.sqlite", "select * from t"))
```



```
1 cd where you can load smart-gnuplotter
2 pydoc -p port smart-gnuplotter
```

and open `localhost:port` with your browser

# Are You Brainwashed?



## Introduction

## Basics

## Walking through a real example

## Settings

### Serial performance

Smarter parameterization

GFLOPS with cores

Speedup

## Overlaying ideal speedup

Confidence interval

## Notes on results





# A real example : graphs we want to show

Let's say we would like to show three graphs:

- ▶ **serial** : compares serial performance among program types, for each architecture
- ▶ **gflops** : shows GFLOPS with cores, for each architecture and program type
- ▶ **speedup** : shows speedup with cores, for each architecture and program type

- ```
1 g.graphs(("matrix.sqlite",
2         r'''select type,avg(gflops_per_sec) from a
3 where ppn=1 and arch="%(arch)s" and M=%(M)s
4 group by type'''),
5         arch=["barcelona", "nehalem", "sandybridge"],
6         M=[704, 1088])
```

## do\_sql method for a smarter parameterization

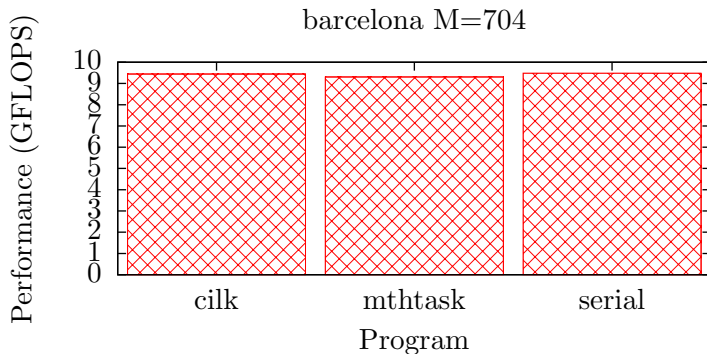
- ▶ You often want to ask database to determine parameters
- ▶ do\_sql method just does that

```


1 Ms=g.do_sql(db,"select distinct M from a where M > 500",single_col=1)
2 archs=g.do_sql(db,"select distinct arch from a",single_col=1)
3 g.graphs(("matrix.sqlite",
4           r'''select type,avg(gflops_per_sec) from a
5 where ppn=1 and arch="% (arch)s" and M=% (M)s
6 group by type'''),
7           arch=archs, M=Ms, graph_vars=["arch", "M"])

```


# Barcelona, $M = 704$



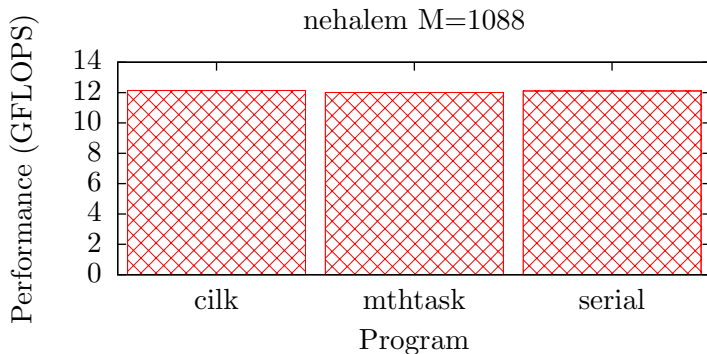
# Barcelona, $M = 1088$



# Nehalem, $M = 704$




# Barcelona, $M = 1088$






## Sandy Bridge, $M = 704$



# Sandy Bridge, $M = 1088$



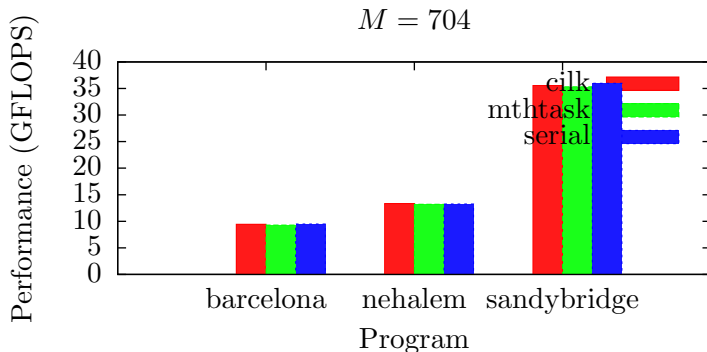
# The real code that saves the 6 graphs

```

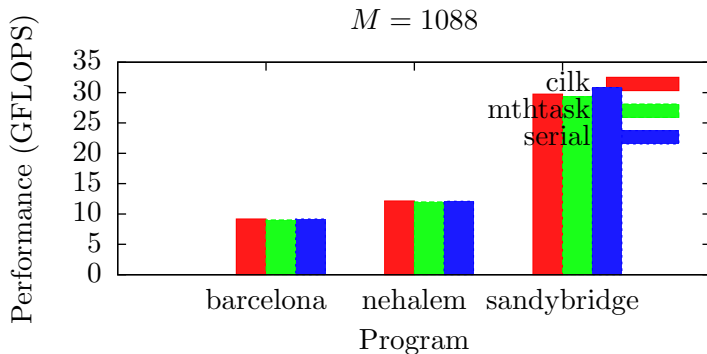
1 query = r'''select type,avg(gflops_per_sec) from a
2 where M = %(M)s and ppn = 1 and arch = "%(arch)s"
3 group by type'''
4
5 g.graphs(("matrix.sqlite", query),
6         output="graphs/serial_%(arch)s_%(M)s.tex",
7         plot_title="",
8         plot_with="boxes fs pattern 1",
9         boxwidth="0.9 relative",
10        graph_title="%(arch)s M=%(M)s",
11        yrange="[0:]",
12        xlabel="Program",
13        ylabel="Performance (GFLOPS)",
14        M=Ms, arch=archs, graph_vars=[ "arch", "M" ])

```

Putting three in a single graph with histogram :  
 $M = 704$



$M = 1088$



## The real code that saves the 2 graphs

```


1 query = r'''select arch,avg(gflops_per_sec) from a
2 where M = %(M)s and ppn = 1 and type = "%(typ)s"
3 group by arch
4 order by arch'''
5 g.graphs(("matrix.sqlite", query),
6         output="graphs/serial_%(M)s",
7         graph_title="$M=%(M)s$",
8         plot_title="% (typ)s",
9         plot.with="histogram fs solid 0.9",
10        using="2",
11        yrange="[0:]",
12        xlabel="Program",
13        ylabel="Performance (GFLOPS)",
14        M=Ms, typ=all_types, graph_vars=[ "M" ])

```

## GFLOPS with cores

```
1 select ppn,avg(gflops_per_sec) from a
2 where type="%(typ)s" and M="%(M)s and arch="%(arch)s"
3 group by ppn
```

Note: I didn't bind workers to cores for Cilk









# Sandy Bridge, $M = 704$

Note: TurboBoost was on




# Sandy Bridge, $M = 1088$

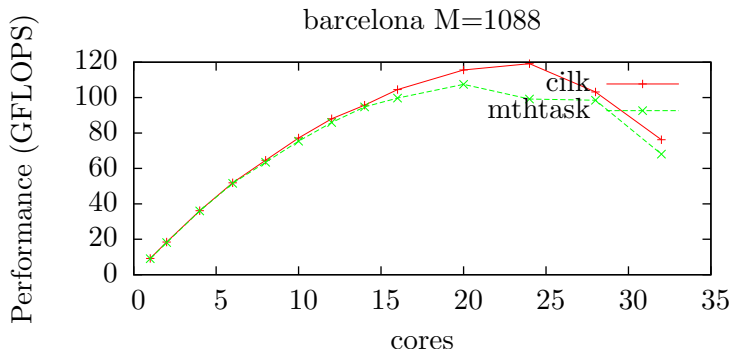
Note: TurboBoost was on



# Barcelona, $M = 704$



# Barcelona, $M = 1088$



The real code that shows the 6 graphs and save them

```

1 Ms = g.do_sql(db, "select distinct M from a where M > 500",
    single_col=1)
2 archs = g.do_sql(db, "select distinct arch from a", single_col=1)
3 para_types = g.do_sql(db, "select distinct type from a where ppn > 1",
    single_col=1)
4
5
6 query = r'''select ppn,avg(gflops_per_sec) from a
7 where type="%(typ)s" and M="%(M)s" and arch="%(arch)s"
8 group by ppn'''
9
10 g.graphs(("matrix.sqlite", query),
11         output="graphs/gflops_%(arch)s_%(M)s.tex",
12         plot_title="%(typ)s",
13         plot_with="linespoints",
14         graph_title="%(arch)s M="%(M)s",
15         xrange=" [0:]",
16         yrange=" [0:]",
17         xlabel="cores",
18         ylabel="GFLOPS",
19         typ=para_types, M=Ms, arch=archs, graph_vars=[ "arch", "M" ])

```

## Translating it into speedup

- Speedup is:

$$\frac{\text{GFLOPS of a program}}{\text{GFLOPS of the serial program for the same parameter}}$$

- So the job is to augment the table with a column of the “GFLOPS of the serial program for the same parameter,”

arch	type	ppn	M	gflops_per_sec	serial_gflops_per_sec
nehalem	serial	1	704	50	50
nehalem	mthtask	1	704	48	50
nehalem	cilk	1	704	49	50
sandybridge	serial	1	704	90	90
sandybridge	mthtask	1	704	88	90
sandybridge	cilk	1	704	89	90
...	...	...	...	...	...

# Preprocessing

```

1 create temp table serial as
2 select arch,M,avg(gflops_per_sec) as serial_gflops_per_sec
3 from a where type = "serial"
4 group by arch,M;
5
6 create temp table b as select * from serial natural join a;


```

Note: you'd better create *temporary* tables only, so you may freely repeat it

## The query

The query itself is easy:


```
1 select ppn, avg(gflops_per_sec / serial_gflops_per_sec) from b
2 where type="% (typ)s" and M=%(M)s and arch="% (arch)s" group by ppn
```









# Sandy Bridge, $M = 704$



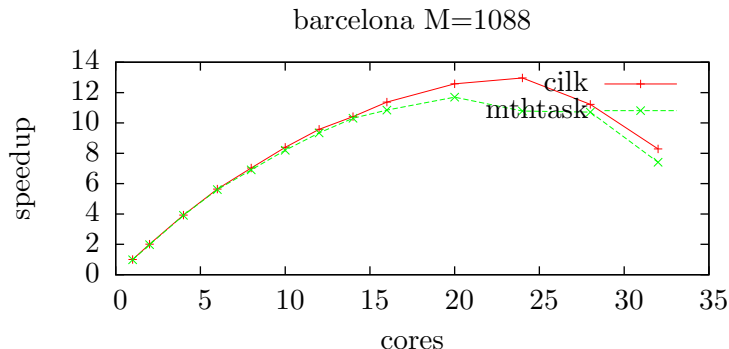
# Sandy Bridge, $M = 1088$



# Barcelona, $M = 704$



# Barcelona, $M = 1088$



## The real code that shows the 6 graphs and save them

```

1  init = r'''create temp table serial as
2  select arch,M,avg(gflops_per_sec) serial_gflops_per_sec
3  from a where type = "serial" group by arch,M;
4
5  create temp table b as select * from serial natural join a;
6  '''
7  query=r'''select ppn,avg(gflops_per_sec / serial_gflops_per_sec) from b
8  where type="% (typ)s" and M=%(M)s and arch="% (arch)s" group by ppn'''
9
10 g.graphs(("matrix.sqlite", query, init),
11          output="graphs/speedup_%(arch)s_%(M)s.tex",
12          plot_title="% (typ)s",
13          plot_with="linespoints",
14          graph_title="% (arch)s M=%(M)s",
15          xrange=" [0:]",
16          yrange=" [0:]",
17          xlabel="cores",
18          ylabel="speedup",
19          typ=para_types, M=Ms, arch=archs, graph_vars=[ "arch", "M" ])

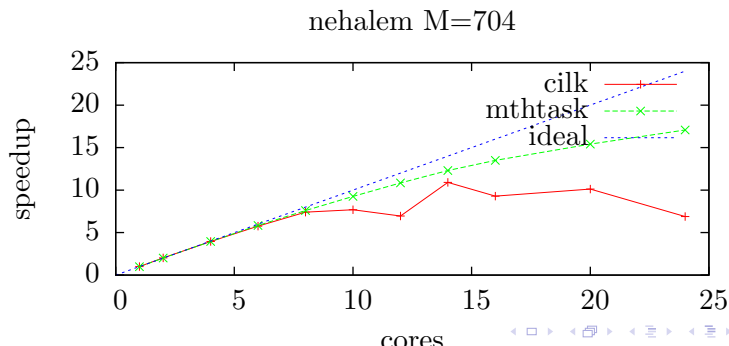
```

## Overlaying “ideal” speedup

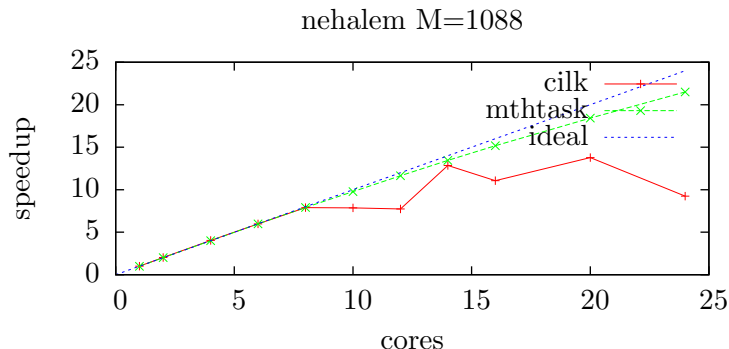
```

1 g.graphs(("matrix.sqlite", query, init),
2         output="graphs/speedup_%(arch)s_%(M)s.tex",
3         overlays=[("x", "plot_title" : "ideal" )],
4         ...)

```




# Nehalem, $M = 1088$







# Sandy Bridge, $M = 704$




# Sandy Bridge, $M = 1088$



# Barcelona, $M = 704$




# Barcelona, $M = 1088$







# Nehalem, $M = 704$



# Nehalem, $M = 1088$




# SandyBridge, $M = 704$







# SandyBridge, $M = 1088$



# Barcelona, $M = 704$



# Barcelona, $M = 1088$



- ▶ **Do not trust Cilk results.** It will become significantly better by binding workers to cores
  - ▶ Now I proved the importance of automating the process!
- ▶ Results on Sandy Bridge (hongo600) will have been affected by TurboBoost, which boosts performance on a single core
  - ▶ We should confirm it by turning TurboBoost off on hongo6xx, but utilizing these dynamic behaviors increasingly sophisticated may be an interesting direction
- ▶ Results on Barcelona is miserable and investigation will lead to a better work stealer

Introduction

Basics

Walking through a real example

Misc.

How you get started (if brainwashed)?



## Are You Brainwashed?



Introduction

Basics

Walking through a real example

Misc.

How you get started (if brainwashed)?



An exclusive offer for you!

- I put a template directory so you can get started now!

```
1 cp -r labdocs/common/tau/new_doc your_new_document
2 cd your_new_document
3 make
```

- It has a structure putting figures, images, databases, python scripts, and Makefile to compile everything into pdf.

## The document template

- ▶ Just 'make' and you get two pdf files
  - ▶ a document
  - ▶ a beamer slide
- ▶ It uses the following programs and T<sub>E</sub>X packages.
  - ▶ `platex` .tex → .dvi
  - ▶ `dvipdfmx` .dvi → .pdf
  - ▶ `inkscape` .svg → .eps
  - ▶ `convert` and `ebb` to get bounding box of them
  - ▶ `convert` .jpg/.png/.gif → .eps
  - ▶ `graphicx`, `listings`, `dvipdfm` T<sub>E</sub>X packages
  - ▶ `beamer` T<sub>E</sub>X package (for slide)

## Other things that will be useful

- ▶ generate  $\text{\TeX}$  table (should be easy)
- ▶ generate Excel file (useful when you prefer/need PowerPoint)