

# How to make T<sub>E</sub>X files with tons of graphs (part II)

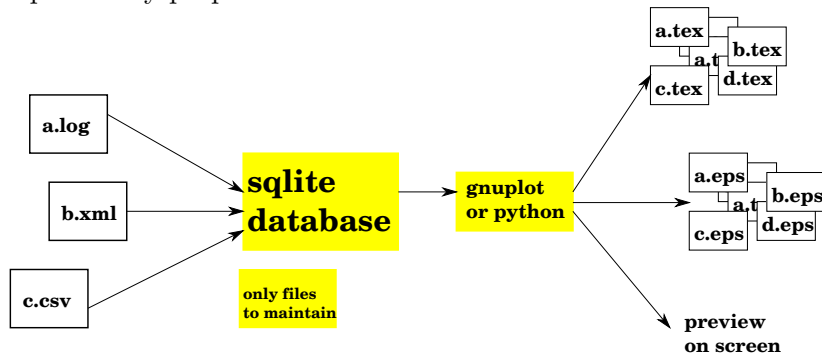


# Context


- ▶ You want to create a T<sub>E</sub>X file with lots of graphs easily
- ▶ You want to *automate the entire process*, from running experiments to producing T<sub>E</sub>X document with graphs, so you can *painlessly repeat the experiment* and update data in the document
- ▶ With ad-hoc solutions, your directory easily screws up with too many data files and scripts you never understand a week later

# The practice

I previously proposed:



# The practice




- ▶ You maintain only 2 files for all graphs
  - ▶ an sqlite3 database that has *all raw data*
  - ▶ a gnuplot or a script file that generates data to gnuplot
- ▶ I talked about a command, `txt2sql`, which make it straightforward to convert text files (log) to sqlite3 database
- ▶ *Generating lots of graphs from a database was still painful, and I now address it*

# What I made this time?

- ▶ A small python library to interface with gnuplot command
- ▶ There is python-gnuplot package, but I do not rely on it
- ▶ A tentative name: `smart_gnuplotter.py`

## A simplest example

```
1 import smart_gnuplotter
2 g = smart_gnuplotter.smart_gnuplotter()
3 g.graphs("x*x")
```



## Pausing behavior

- ▶ It generates “pause -1” so gnuplot waits until you enter a newline
- ▶ smart\_gnuplotter then asks what to do for the following graphs:

```
1 's' to suppress future prompts, 'q' to quit, else to continue [s/q/  
  other]?
```


- ▶ You may change it by:

```
1 g.default_pause = 0
```

# Writing multiple curves in a single graph

```
1 g.graphs("%(a)s * x", a=[1,2,3])
```


- ▶ The basic form is to parameterize the expression with *%(var)s* and supply its values with *var=list*






# More curves in a single graph ...

```
1 g.graphs("%(a)s*x*x+%(b)s*x+%(c)", a=[1,2,3], b=[4,5,6], c=[7,8,9])
```




# Writing multiple graphs in a single shot (1)


```
1 g.graphs("%(a)s*x*x+%(b)s*x+%(c)", a=[1,2,3], b=[4,5,6], c=[7,8,9],
2      graph_vars=["a"])
```



## Writing multiple graphs in a single shot (2)



## Writing multiple graphs in a single shot (3)



## Parameters that change together

- ▶ So the basic form is:

```
1 g.graphs("expression", var=values, var=values, ...)
```

- ▶ But you might not want to generate all combinations
- ▶ For example, you may want to have  
 $(a, b) = ("hoge", 10), ("bar", 20)$ , not  $("hoge", 20)$  or  $("bar", 10)$
- ▶ You may do so by having a parameter assuming tuple values:

```
1 X=[("hoge", 10), ("bar", 20)]
```

and refer to them as  $\%(X[0])s$ ,  $\%(X[1])s$ , etc.

## Looking what's going on

```
1 g.graphs(..., gpl_file="hoge.gpl")
```


and look at hoge.gpl:

```
1 set terminal wxt
2 plot x*x
3 pause -1
```

- ▶ Specify `gpl_file` and you will get the file given to gnuplot
- ▶ When something went wrong, looking at this file is often the simplest way to diagnose

## Specifying various attributes

```
1 g.graphs("%(a)s*x*x+% (b)s*x+% (c)s", a=[1,2,3], b=[4,5,6], c=[7,8,9],
2       graph_vars=["a"], graph_attr=r'''
3 set title "%(a)sx^2+bx+c$"
4 ''',
5       curve_attr=r'''title "%(a)sx^2+% (b)sx+% (c)s$"'''')
```



## Specifying various attributes

- ▶ `graph_attr` is whatever comes before the 'plot' command
- ▶ `curve_attr` is whatever comes after each curve

```
1 g.graphs(E,  
2     graph_attr=graph_attr,  
3     curve_attr=curve_attr)
```

```
1 graph_attr  
2 plot E curve_attr, E curve_attr, ...
```



## Shortcut for frequently used attributes

Frequently used attributes can be directly set by keyword arguments

```
1 g.graphs(..., var=..., var=..., ...)
```

- ▶ Graph attributes
  - ▶ graph\_title
  - ▶ terminal
  - ▶ output
  - ▶ xrange, yrange
  - ▶ xlabel, ylabel
  - ▶ boxwidth
- ▶ Curve attributes
  - ▶ curve\_title
  - ▶ curve\_with

## Other things to plot

- ▶ As it is normal in gnuplot, you may plot

- ▶ datafile

```
1 g.graphs('filename', ...)
```

```
2
```

- ▶ output of a command

```
1 g.graphs('>cmd', ...)
```

```
2
```

- ▶ Besides, you may plot

- ▶ data in python list

```
1 g.graphs([(1,2),(3,4),...], ...)
```


```
2
```

- ▶ output of an SQLite query (later)

## Bonus : writing graphs with symbolic $x$ -axis

- ▶ It's tedious to show a graph with symbolic  $x$ -axis
- ▶ `smart_gnuplotter` does all the work you need to do for it

```
1 g.graphs([("a",10), ("b",20), ("c", 30)],  
2         xrange="0:",  
3         boxwidth="0.9 relative",  
4         curve_with="boxes fs pattern 1")
```



## Interacting with sqlite3: where the real power comes from

```
1 g.graphs((database, query, init_string, init_file, udfs, udas,  
           udfs), ...)
```

- ▶ When you give a tuple as the first argument, a database query is executed and the result treated as data
- ▶ `init_string`, `init_file`, `udfs`, `udas`, `udcs` are optional
- ▶ The above  $\approx$

```
1 co = sqlite3.connect(database)  
2 for name,arity,f in udfs:  
3     co.create_functions(name,arity,f)  
4 for name,arity,f in udas:  
5     co.create_aggregates(name,arity,f)  
6 for name,f in udfs:  
7     co.create_collations(name,f)  
8 co.execscript(init_string)  
9 co.execscript(content of init_file)  
10 co.execute(query)
```


## A basic example with sqlite3

Say a database example.sqlite contains the following table

```
1 select * from t
2 0|1
3 2|3
4 4|5
```

Then

```
1 g.graphs(("example.sqlite", "select * from t"))
```



## A real example (1)

Database 'a' contains all results from matrix multiply

```
1 sqlite> .schema  
2 CREATE TABLE a (arch, type, ppn, M, gflops_per_sec, ...);
```

- ▶ arch : architecture (barcelona, nehalem, sandybridge)
- ▶ type : program type (serial, MassiveThreads, Cilk)
- ▶ ppn : number of cores (1, 2, 4, 6, ...)
- ▶ M : matrix size
- ▶ gflops\_per\_sec : performance (GFLOPS)

Experiments are repeated, so there are many date of the same (arch, type, ppn, M)

## A real example (2)

Let's say we would like to show three graphs:

- ▶ serial: compares serial performance among program types, for each architecture
- ▶ gflops: shows GFLOPS with cores, for each architecture and program type
- ▶ speedup: shows speedup with cores, for each architecture and program type

# Serial

- ▶ STEP 1: take some time interacting with sqlite3 to come up with a right query showing data for a single graph

```
1 select type,avg(gflops_per_sec) from a
2 where ppn=1 and arch="nehalem" and M=704
3 group by type
```

- ▶ STEP 2: parameterize it:

```
1 g.graphs(("matrix.sqlite",
2         r'''select type,avg(gflops_per_sec) from a
3 where ppn=1 and arch="% (arch)s" and M=% (M)s
4 group by type'''),
5         arch=["barcelona", "nehalem", "sandybridge"],
6         M=[704, 1088])
```




## do\_sql method for a smarter parameterization

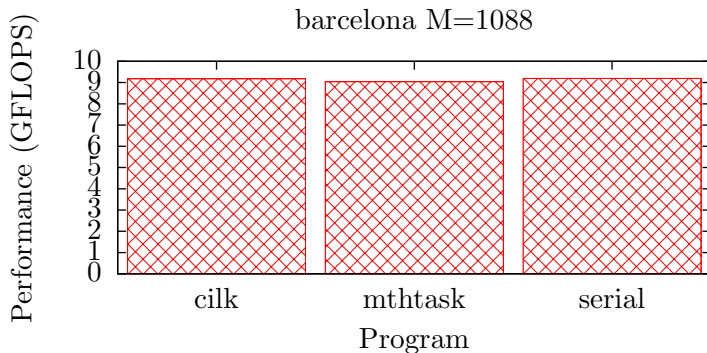
- ▶ You often want to ask database to determine parameters
- ▶ do\_sql method just does that

```
1 Ms = g.do_sql(db, "select distinct M from a where M > 500", single_col
    =1)
2 archs = archs = g.do_sql(db, "select distinct arch from a", single_col
    =1)
3 g.graphs(("matrix.sqlite",
4         r'''select type,avg(gflops_per_sec) from a
5 where ppn=1 and arch="%(arch)s" and M=%(M)s
6 group by type'''),
7         arch=archs, M=Ms)
```


## Barcelona, $M = 704$




Barcelona,  $M = 1088$




Nehalem,  $M = 704$



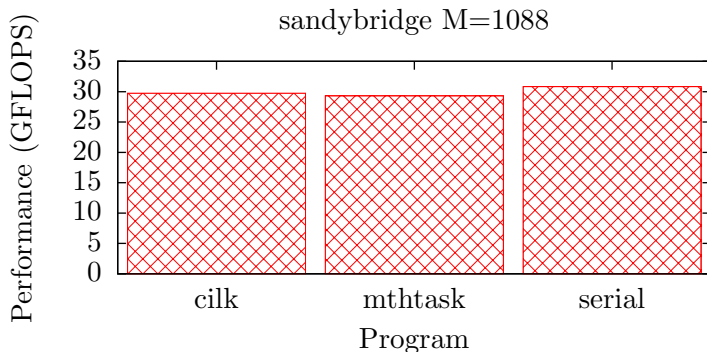
Barcelona,  $M = 1088$



## Sandy Bridge, $M = 704$



## Sandy Bridge, $M = 1088$



## The real code that saves the 6 graphs


```
1 query = r'''select type,avg(gflops_per_sec) from a
2 where M = %(M)s and ppn = 1 and arch = "%(arch)s"
3 group by type'''
4
5 g.graphs((db, query),
6          output="graphs/serial_%(arch)s_%(M)s.tex",
7          curve_title="",
8          curve_with="boxes fs pattern 1",
9          boxwidth="0.9 relative",
10         graph_title="%(arch)s M=%(M)s",
11         yrange="[0:]",
12         xlabel="Program",
13         ylabel="Performance (GFLOPS)",
14         M=Ms, arch=archs, graph_vars=[ "arch", "M" ])
```




# GFLOPS with cores

```
1 select ppn,avg(gflops_per_sec) from a
2 where type="%(typ)s" and M=%(M)s and arch="%(arch)s"
3 group by ppn
```


Note: I didn't bind workers to cores for Cilk



## Nehalem, $M = 1088$




## Sandy Bridge, $M = 704$




## Sandy Bridge, $M = 1088$

Note: TurboBoost is probably on




## Barcelona, $M = 704$

Note: TurboBoost is probably on



## Barcelona, $M = 1088$



## The real code that shows the 6 graphs and save them

```
1 Ms      = g.do_sql(db, "select distinct M from a where M > 500",
    single_col=1)
2 archs = g.do_sql(db, "select distinct arch from a", single_col=1)
3 para_types = g.do_sql(db, "select distinct type from a where ppn > 1",
    single_col=1)
4
5
6 query = r'''select ppn,avg(gflops_per_sec) from a
7 where type="%(typ)s" and M="%(M)s" and arch="%(arch)s"
8 group by ppn'''
9
10 g.graphs(("matrix.sqlite", query),
11         output="graphs/gflops_%(arch)s_%(M)s.tex",
12         curve_title="%(typ)s",
13         curve_with="linespoints",
14         graph_title="%(arch)s M="%(M)s",
15         xrange=" [0:]",
16         yrange=" [0:]",
17         xlabel="cores",
18         ylabel="GFLOPS",
19         typ=para_types, M=Ms, arch=archs, graph_vars=[ "arch", "M" ])
```

# Translating it into speedup

- ▶ Speedup is:

$$\frac{\text{GFLOPS of a program}}{\text{GFLOPS of the serial program for the same parameter}}$$

- ▶ So the job is to augment the table with a column of the “GFLOPS of the serial program for the same parameter,”

arch	type	ppn	M	gflops_per_sec	serial_gflops_per_sec
nehalem	serial	1	704	50	50
nehalem	mthtask	1	704	48	50
nehalem	cilk	1	704	49	50
sandybridge	serial	1	704	90	90
sandybridge	mthtask	1	704	88	90
sandybridge	cilk	1	704	89	90
...	...	...	...	...	...



# Preprocessing


```
1 create temp table serial as
2 select arch,M,avg(gflops_per_sec) serial_gflops_per_sec
3 from a where type = "serial"
4 group by arch,M;
5
6 create temp table b as select * from serial natural join a;
```

Note: you'd better create *temporary* tables only, so you may freely repeat it


# The query

The query itself is easy:


```
1 select ppn, avg(gflops_per_sec / serial_gflops_per_sec) from b
2 where type="% (typ)s" and M=% (M)s and arch="% (arch)s" group by ppn
```



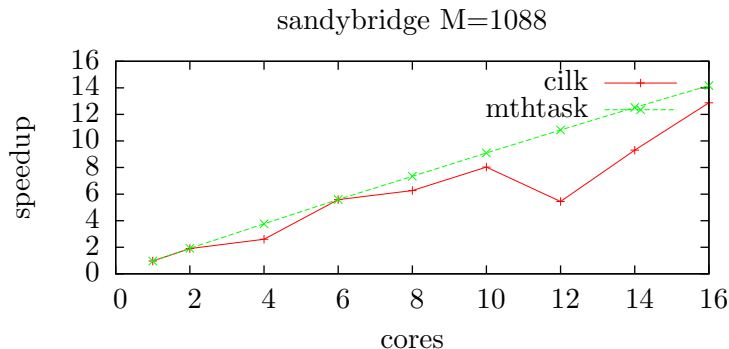
# Nehalem, $M = 1088$



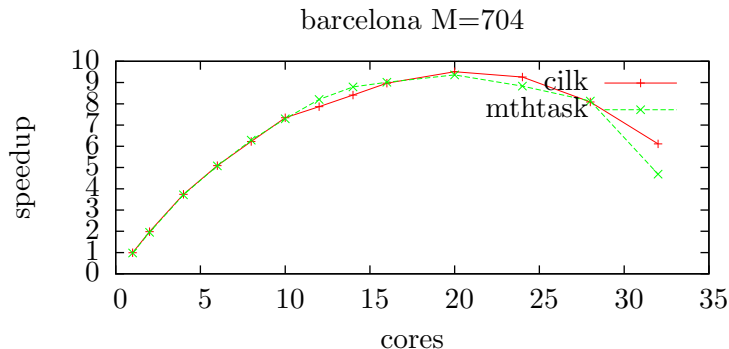
## Sandy Bridge, $M = 704$




## Sandy Bridge, $M = 1088$



## Barcelona, $M = 704$



## Barcelona, $M = 1088$




## The real code that shows the 6 graphs and save them

```
1 init = r'''create temp table serial as
2 select arch,M,avg(gflops_per_sec) serial_gflops_per_sec
3 from a where type = "serial"
4 group by arch,M;
5
6 create temp table b as select * from serial natural join a;
7 '''
8 query = r'''select ppn,avg(gflops_per_sec / serial_gflops_per_sec) from
9         b
10        where type="%(typ)s" and M=%(M)s and arch="%(arch)s" group by ppn'''
11 g.graphs(("matrix.sqlite", query, init),
12         output="graphs/speedup_%(arch)s_%(M)s.tex",
13         curve_title="%(typ)s",
14         curve_with="linespoints",
15         graph_title="%(arch)s M=%(M)s",
16         xrange=" [0:]",
17         yrange=" [0:]",
18         xlabel="cores",
19         ylabel="speedup",
20         typ=para.types, M=Ms, arch=archs, graph_vars=["arch", "M"] )
```




# Overlaying “ideal” speedup


```
1 g.graphs("matrix.sqlite", query, init),  
2     output="graphs/speedup_%(arch)s_%(M)s.tex",  
3     overlays=[("x", 'title "ideal"')],  
4     ...)
```




## Nehalem, $M = 1088$




## Sandy Bridge, $M = 704$




## Sandy Bridge, $M = 1088$



## Barcelona, $M = 704$



## Barcelona, $M = 1088$



## Note on the results

- ▶ Do not trust Cilk results. It will become significantly better by binding workers to cores
- ▶ Now I proved the importance of automating the process!
- ▶ Results on Sandy Bridge (hongo600) will have been affected by TurboBoost, which boosts performance on a single core
- ▶ We'd better to confirm it by turning TurboBoost off on hongo6xx, but utilizing these dynamic behaviors increasingly sophisticated may be an interesting direction
- ▶ Results on Barcelona need true serious investigation and will be a research theme for us

# Questions? Objections?

